



University of Trieste
Data Management for Big Data Course
Academic Year 2022–2023

Data Warehouse case study

Davide Capone* Sandro Junior Della Rovere† Enrico Stefanel‡

Last update on 2023-06-21T13:32:43Z.

1 Introduction

The aim of this project is to study an efficient implementation of a Postgres database using the public TPC-H benchmark, which can be considered as a Big Data source.

1.1 TPC-H benchmark database

The TPC-H benchmark can be downloaded

1.1.1 Database statistics

The benchmark is composed by the following tables:

- CUSTOMER, with 1 500 000 tuples (312 MB);
- LINEITEM, with 59 986 052 tuples (11 GB); the main attributes that are going to be used are:
 - `l_extendedprice` (1 351 462 distinct values, i.e. there is an average of 44 tuples with the same value, that range from 900.91 to 104 949.50),

*davide.capone@studenti.units.it

†sandrojunior.dellarovere@studenti.units.it

‡enrico.stefanel@studenti.units.it

- `l_discount` (11 distinct values, i.e. there is an average of 5 453 277 tuples with the same value, that range from 0.00 to 0.10),
 - `l_returnflag` (which can assume values `A`→not returned, `R`→returned, `N`→not yet delivered),
 - `l_commitdate`,
 - `l_receiptdate`,
- NATION, with 25 tuples (24kB);
- ORDERS, with 1 500 000 tuples (2481 kB); the main attributes that are going to be used are:
 - `o_orderdate`,
- PART, with 2 000 000 tuples (363 MB); the main attributes that are going to be used are:
 - `p_type`;
- PARTSUPP, with 8 000 000 tuples (1535 MB);
- REGION, with 5 tuples (24kB);
- SUPPLIER, with 100 000 tuples (20 MB).

Other attributes have been used (e.g., keys used to join relations), but statistics about them have been omitted for lack of usefulness.

2 Export/import revenue value

2.1 Naïve Implementation

```

1  WITH query1 AS (
2  SELECT
3      EXTRACT (YEAR FROM o_orderdate) AS _year,
4      EXTRACT (QUARTER FROM o_orderdate) AS _quarter,
5      EXTRACT (MONTH FROM o_orderdate) AS _month,
6      c_regionname,
7      c_nationname,
8      c_name,
9      s_regionname,
10     s_nationname,
11     s_name,
12     p_type,
13     SUM(l_extendedprice * (1 - l_discount)) AS revenue
14 FROM lineitem_orders

```

```

15     JOIN part ON l_partkey = p_partkey
16     JOIN supplier_location ON (s_suppkey = l_suppkey)
17     JOIN customer_location ON (c_custkey = o_custkey)
18 WHERE s_nationkey <> c_nationkey
19 GROUP BY
20     _year,
21     _quarter,
22     _month,
23     c_regionkey,
24     c_regionname,
25     c_nationkey,
26     c_nationname,
27     c_custkey,
28     c_name,
29     s_regionkey,
30     s_regionname,
31     s_nationkey,
32     s_nationname,
33     s_suppkey,
34     s_name,
35     p_type
36 )
37 SELECT * FROM query1;

```

3 Late delivery

It is asked to retrieve the number of orders where at least one “lineitem” has been received later than the committed date. The aggregation should be performed with the Month → Year roll-up, and the (Customer’s) Nation → Region roll-up.

Lorem ipsum. . .

4 Returned item loss

It is asked to retrieve the *revenue loss* for customers who might be having problems with the parts that are shipped to them, where a *revenue lost* is defined as $\text{SUM}(l_extendedprice * (1 - l_discount))$ for all qualifying “lineitems”.

4.1 Naïve implementation

```

1 WITH lineitem_orders AS (

```

```

2      SELECT
3          o_orderkey,
4          l_partkey,
5          l_suppkey,
6          o_orderdate,
7          o_custkey,
8          l_extendedprice,
9          l_discount,
10         l_returnflag,
11         l_commitdate,
12         l_receiptdate
13     FROM lineitem JOIN orders ON (l_orderkey = o_orderkey)
14 ),
15 query3 AS (
16     SELECT
17         EXTRACT (YEAR FROM o_orderdate) AS _year,
18         EXTRACT (QUARTER FROM o_orderdate) AS _quarter,
19         EXTRACT (MONTH FROM o_orderdate) AS _month,
20         c_name,
21         SUM(l_extendedprice * (1 - l_discount)) AS returnloss
22     FROM
23         lineitem_orders
24         JOIN customer ON o_custkey = c_custkey
25     WHERE
26         l_returnflag = 'R'
27         -- AND c_name = 'Customer#000129976'
28         -- AND EXTRACT (QUARTER FROM o_orderdate) = 1
29     GROUP BY
30         _year,
31         _quarter,
32         _month,
33         c_custkey,
34         c_name
35 )
36 SELECT * FROM query3;

```